

Team Delft's Robot Winner of the Amazon Picking Challenge 2016

Carlos Hernandez^{1(✉)}, Mukunda Bharatheesha¹, Wilson Ko², Hans Gaiser²,
Jethro Tan¹, Kanter van Deurzen², Maarten de Vries², Bas Van Mil²,
Jeff van Egmond¹, Ruben Burger¹, Mihai Morariu², Jihong Ju¹,
Xander Germann¹, Ronald Ensing², Jan Van Frankenhuyzen¹,
and Martijn Wisse^{1,2}

¹ Robotics Institute, Delft University of Technology,
Mekelweg 2, 2628 Delft, CD, The Netherlands
c.h.corbato@tudelft.nl

² Delft Robotics, B.V., Mijnbouwstraat 120,
2628 Delft, RX, The Netherlands

Abstract. This paper describes Team Delft's robot, which won the Amazon Picking Challenge 2016, including both the Picking and the Stowing competitions. The goal of the challenge is to automate pick and place operations in unstructured environments, specifically the shelves in an Amazon warehouse. Team Delft's robot is based on an industrial robot arm, 3D cameras and a customized gripper. The robot's software uses ROS to integrate off-the-shelf components and modules developed specifically for the competition, implementing Deep Learning and other AI techniques for object recognition and pose estimation, grasp planning and motion planning. This paper describes the main components in the system, and discusses its performance and results at the Amazon Picking Challenge 2016 finals.

Keywords: Robotic system · Warehouse automation · Motion planning · Grasping · Deep learning

1 Introduction

The Amazon Picking Challenge (APC) was launched by Amazon Robotics in 2015 [3] to promote research into robotic manipulation for picking and stocking of products. These tasks are representative of the current challenges that warehouse automation faces nowadays. The unstructured environment and the diversity of products require new robotic solutions. Smart mechanical designs and advanced artificial intelligence techniques need to be combined to address the challenges in object recognition, grasping, dexterous manipulation or motion planning.

Amazon chose 16 teams from all over the world to participate in the finals at RoboCup 2016. Team Delft won both the picking and the stowing challenges. Section 2 discusses Team Delft's approach, explaining its design principles and

the robot hardware. Section 3 details the robot control and all the components integrated for object detection, grasp and motion planning. Finally Sects. 4 and 5 discuss the competition results and the lessons learned. The purpose of this paper is to provide a comprehensive analysis of the complete development of an advanced robotic system that has to perform in real-world circumstances.

2 The Amazon Picking Challenge 2016

The Amazon Picking Challenge 2016 included two competitions: in the Picking Task 12 items from the competition product set had to be picked from an Amazon shelving unit and placed in a tote; in the Stowing Task it was the other way around: 12 items were to be picked from the tote and stowed into the shelf. The maximum allotted time to fulfil each task was 15 min and the system had to operate autonomously. A file containing the task order was given to the system, which included the initial contents of the shelf's bins and the tote, and it had to produce a resulting file indicating the location of all the products.

The set of 39 items used in the challenge (Fig. 1) were representative of those in an Amazon warehouse. Books, cubic boxes, clothing, soft objects, and irregularly shaped objects represented realistic challenges such as reflective packaging, different sizes or deformable shapes. The items could be placed in any orientation inside the bins, sometimes cluttering them, and the target product could be partially occluded by others.

Teams had to place their robots in a $2\text{ m} \times 2\text{ m}$ workcell, no closer than 10 cm from the shelf. The workspace also posed important challenges to perception and manipulation. The shelf was a metal and cardboard structure divided into a matrix of 3 by 4 bins. The bins were narrow but deep, which limited the manoeuvrability inside and required a long reach. Additionally, the shelf construction resulted in significant deviations in reality from its ideal geometric model.

The performance of the robots during the picking and the stowing tasks was evaluated by giving points for correctly placed items and subtracting penalty points for dropping, damaging or misplacing items. A correct operation could receive 10, 15 or 20 points depending on the cluttering of the bin. Additional bonus points were given for specially difficult objects, for maximum scoring of 185 points in the Picking Task and 246 points in the Stowing Task.

3 Team Delft's Robot

Team Delft was a joint effort of the Robotics Institute of the Delft University of Technology [11] and the robot integrator company Delft Robotics B.V. [4] to participate in the APC 2016. Amongst TUD Robotics Institute research lines is the development of flexible robots capable of automatizing small-scale productions, simplifying their installation and reconfiguration, e.g. through automatic calibration or online self-generated motions. Delft Robotics is a novel systems integrator making these new robotic technologies available to all kind of manufacturing companies. Both parties are closely collaborating within the




Fig. 1. The 39 items in the Amazon Picking Challenge 2016

Factory-in-a-day EU project [5] to reduce installation time and cost of robot automation. Team Delft’s goal was to demonstrate and validate this approach in such a challenging industrial benchmark as the APC. The team did not adapt and tune an extant pick-and-place solution to participate in APC, but developed the best solution possible with extant industrial hardware and as many off-the-shelf software components as possible. For that the robot control was based on the ROS framework [7].

The remaining of this section describes the main ideas behind Team Delft's robotic solution.

3.1 Robot Concept

The team analysed the results of the previous edition of the Amazon Picking Challenge in 2015 [3], and decided that making the system *robust* and *fast* was key to win. These characteristics allow the system to perform several attempts to pick each target item, and move occluding objects around if necessary. We also learned that suction was the better performing grasp option, which we confirmed in early tests.

The solution designed is based on an industrial robot arm, a custom made gripper and 3D cameras, as shown in Fig. 2. For the robot arm we chose a 7 degrees of freedom SIA20F Motoman mounted on an horizontal rail perpendicular to the shelf. The resulting 8 degrees of freedom allowed the system to reach all the bins with enough manoeuvrability to pick the target objects.




Fig. 2. Team Delft robot setup in the APC workcell.

We customized our own gripper to handle all the products in the competition (see Fig. 3). It has a lean footprint to manoeuvre inside the bins, and a 40 cm length to reach objects at the back. It includes a high flow suction cup at the end, with a 90° rotation allowing two orientations, and a pinch mechanism for the products difficult to suck. Both the suction cup rotation and the pinch mechanism are pneumatically actuated. A vacuum sensor provides boolean feedback whether the suction cup holds anything. For object detection a 3D camera is mounted in the gripper to scan the bins, while another one is fixed on a pole above the tote.




Fig. 3. Team Delft gripper.

The tote is placed on a frame attached to the robot rail. The compressor and the vacuum pump required to actuate the gripper are mounted on another frame that attached to the rail base, so the whole set up could be easily moved in three big blocks. Robust and easy transportation and installation were important requirements.

3.2 Control Pipeline

The system control is based on the sense-plan-act paradigm and path planning for robot motion. This allows for potentially optimal motions, at the cost of more precise sensing information. First, the task is decomposed into a set of pick and place operations on the target items. Then, for each operation in the Picking task the sense-plan-act cycle proceeds as follows¹. First in the *sense* step the robot moves to take an image of the bin containing the first target item to locate it and get the obstacles information. Then, during the *plan* step a grasping strategy and candidate pose for the gripper to grab it are computed, and a motion plan is generated to approach, grasp and retreat from the bin with the item. Following, in the *act* step the gripper is configured for the selected strategy and the complete motion is executed, including gripper activation to suck or pinch-grasp the item. The vacuum seal in the suction cup is checked to confirm a successful pick. If so, the robot moves to deposit the item in the tote, using simple drop-off motions. This cycle is repeated till all target items are picked. For the Stowing task the loop operates similarly until all items in the tote are stowed in the shelf.

4 Robot Software

Team Delft was fully committed to the ROS-Industrial initiative [9] that aims to create industry-ready, advanced software components to extend the capabilities of factory robots. The robot software is thus based on the ROS framework [7]. We found that the flexibility, modularity and tools provided by ROS allowed us to address the requirements for autonomy and high and reliable performance in the competition, and facilitated development.

The ROS component-based approach allowed for the integration of the different components for task management, object detection, pose estimation, grasping and motion planning into a robust architecture. Following we describe them.

4.1 Task Management

On top of the architecture sits the task manager, responsible for decomposing the Pick and the Stow tasks into a plan of pick and place operations, and manages the state of fulfilment of the whole task. It encodes the competition rules to

¹ A video demonstrating the pipeline can be found here: <https://www.youtube.com/watch?v=PKgFy6VUC-k>.

maximize the scoring, by planning first those operations that scored more points, and keeps track of the location of all the items.

A central coordinator module coordinates the execution of each pick and place operation following the sequential flow presented in Sect. 3.2. It was implemented as a ROS SMACH [2] state machine.

The system can handle some failures applying fallback mechanisms to continue operation. For example, if the robot cannot find the target item, or estimate its pose, it tries different camera viewpoints, then if the problem persists it postpones that target and moves to the next operation. The system can detect if a suction grasp failed by checking the vacuum sealing after execution of the complete grasp and retreat action. If there is no seal the robot assumes the item dropped inside the bin and retries the pick later. If the seal is broken during the placement in the tote, the item is assumed to have dropped in the tote.

4.2 Object Recognition and Pose Estimation

The robot’s pipeline involved detecting the target item within the bin or the tote and obtaining a grasp candidate using an estimation of its pose or its centroid, in the case of deformable items. Difficulties included narrow view angles and poor lighting conditions and reflections inside the shelf.

Firstly, the system acquires the 3D² and RGB image with an Ensenso N35 camera. For that, in the Picking task the robot moves the gripper to a pre-defined location in front of the desired bin. In the Stowing task the image is taken by the camera fixed over tote. Then object detection consists of two main steps:

Object Recognition. First, a deep neural network based on Faster R-CNN [8] classifies the objects in the RGB image and extracts their bounding boxes. A pre-trained neural network was further trained to create the two models used for object recognition in both the picking and the stowing tasks. A dataset of about 20 K images of the products in different orientations and with random backgrounds was created to train a “base” model. Then this model was trained with around 500 labelled images of real occurrences of the products in the shelf and in the tote to generate the final recognition models. The result was an almost flawless detection within 150 ms of all the products present in any bin or tote image, as shown in Fig. 4.

Pose Estimation. Pose estimation of non-deformable products was done using Super 4PCS [6] to match the filtered PointCloud of the target item with a CAD model of the object. The 3D information of the bin or the tote is also used later during motion planning for collision detection. Reflections due to packaging and the difficult lighting conditions inside the bin resulted in scarce and noisy 3D data for some products. This proved a big difficulty for the pose estimation method. We included heuristics to correct estimations, e.g. objects cannot be floating on the bin, and also the mentioned fall-back mechanism to take additional images.

² The 3D data format used was PointCloud.

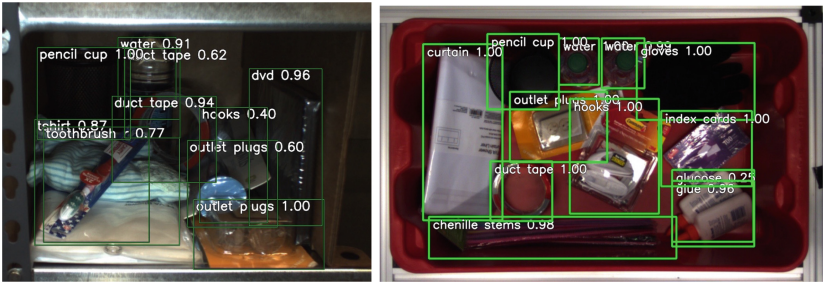


Fig. 4. Example result of the object detection module based on Faster R-CNN for bin and tote images. The estimated bounding boxes are depicted in green and labelled with the identification proposal and the confidence. (Color figure online)

4.3 Grasping and Manipulation

The grasp and manipulation solution is customised to our gripper and our path planning approach. The gripper has three basic modes or configurations (see Fig. 5): *front suction*, *side-top suction*, and *pinch*, each one corresponding to a grasping strategy more suitable different products also depending on the situation.



Fig. 5. The different gripper configurations for the grasping strategies. In the top images, the two configurations for suction rotating the cup. In the image below, the robot is picking the dumbbell using the pinch configuration.

In the *plan* step the best strategy and associated grasp candidate—i.e. a 6D pose to position the gripper—to grasp the target item are chosen, and then the system computes a manipulation plan to move the gripper to the candidate pose, activate it to pick the item, and move out of the bin (or the tote) holding it.

Grasp planning. For non-deformable items pose estimation of the object and offline geometric constraints and user-defined heuristics are used to synthesize grasp candidates. Basically a set of grasp candidates is generated over the surface of the 3D model of the item based on primitive shapes (cylinders, spheres, cones, planes, etc.). These candidates are pruned online using geometry constraints due to the actual item’s estimated pose and gripper limitations, e.g. candidates at the back or the bottom³ of the item are not reachable. Additional heuristics were defined experimentally to prune those grasp candidates specific to each product that proved not suitable. These heuristics were implemented so that new ones could easily be coded for a set of products by including them in the primitive shapes that accounted for different products, while still being able to define ad-hoc constraints that only applied to specific products. Finally, the synthesized candidates are scored based on geometry and dynamic considerations, e.g. poses closer to the centre or mass would tend to provide more stable grasps.

For deformable items the system exploits the power of the suction cup, which is capable of grasping and holding all deformable products in the competition. Instead of computing grasp candidates from the 3D pose estimation of the object, the normals of the segmented object PointCloud are directly used as grasp candidates. The candidates are also scored simply based on the distance to the PointCloud centroid, the closer the better.

Manipulation. Actually grasp planning produces not one but a set of grasp candidates ranked by our scoring criteria. However, the robot might not be able to reach some of these poses with the gripper, due to its kinematic limits or the obstacles, such as the bin or the tote walls, or other items close to the target one. Even if reachable, the robot, with the item attached to the gripper, also needs a retreat trajectory free of obstacles. The first grasp candidate for which a collision-free pick and retreat complete trajectory can be computed is then selected. This will be detailed in Sect. 4.4.

4.4 Motion Planning

For the robot motion strategy we divided the problem considering that the workspace is static (apart from motions due to the robot), and known outside of the shelf. Any online motion planning was required only inside the bins or the tote.

³ These are relative locations assuming we are looking at the object from the tip of the gripper.

Offline Motions. Collision-free trajectories between all relevant locations to approach the bins and the tote and capture the images were computed offline. The trajectories were generated in joint space with the RRT-Connect randomized path planner via MoveIt! [10] and using a URDF⁴ model of the workcell including the shelf, the robot on the rail, the gripper, and all the attached frames and equipment.

Online Cartesian Path Planning. To simplify the manipulation problem inside the bins, only collision-free picks were to be attempted. We defined a cartesian approach based on the MoveIt! pick and place pipeline that took the target grasp candidate and computed a combination of linear segments to *approach*, *contact* grasp the target object, *lift* it after grasping and *retreating* with it. The TRAC-IK library [1] is used for inverse kinematics, configured to enforce minimal configuration changes, and then collision checking is done with MoveIt! using the PointCloud information from the camera.

Robot Motion. This way, for the Picking task, offline motions were used in the *sense* phase to acquire the image of the bin containing the target object and to position the gripper ready to enter the bin. Then, during the *plan* phase the approach, contact, lift and retreat segments were generated online, and a drop-off location chosen and an associated offline trajectory retrieved. Finally, a complete motion plan to pick and place the target item is generated by stitching the cartesian segments and the offline drop-off trajectory. This includes time parametrization and the I/O commands required to configure and activate the gripper for grasping, resulting in a complete trajectory that is executed by the robot in the *act* phase. The MotoROS driver was used and enhanced by Team Delft⁵ to execute the desired trajectories controlling the complete kinematic chain of the robot and the rail, and also the gripper using the robot controller I/O.

5 Competition Results

Team Delft's robot was the champion of the challenge winning both competitions, with an outstanding performance in the Stowing Task⁶. Table 1 shows the final scores for the Amazon Picking Challenge 2016 Pick and Stow competitions. The overall results of the teams improved considerably over the previous APC edition: average scoring for the top 10 teams increased 38% for the Picking Task, specially considering the increased difficulty in this edition, with more cluttered bins. It is also interesting to mention that all the best robots but Team Delft's placed the tote below the shelf, and initially moved a board to act as a ramp so that any items dropping will fall down to the tote. This trick improved scoring.

⁴ Unified Robot Description Format <http://wiki.ros.org/urdf>.

⁵ This contribution, as well as other ROS components developed for APC will be open to the community.

⁶ Video recordings of Team Delft's competition runs can be found here <https://youtu.be/3KlzVWxomqs> (picking) and here <https://youtu.be/AHUuDVdiMfg> (stowing).

We considered this mechanical solution early at the concept brainstorming, but finally discarded it because due to the rail there was no free space for a clean and robust design. We did not want to include any provisional duck-tape solution. However, Team Delft’s robust and fast concept outperformed the rest achieving more successful pick and place operations, which was the aim of the competition.

In the Stowing Task Team Delft’s robot successfully stowed 11 items of the 12, dropping the remaining one while manipulating one of the other products. The system only had to retry one of the picks from the tote, to finish the task in a total time of 7 min 10 s.

Table 1. Amazon Picking Challenge 2016 scores of the best four robots.

	Stowing scores		Picking scores
214	Team Delft	105	Team Delft (0:30 first pick)
186	NimbRo picking	105	PFN (1:07 first pick)
164	MIT	97	NimbRo picking
161	PFN	67	MIT

The picking task proved much harder than the stowing. The robot picked successfully 9 out of 12 items, the first one in only 30 s. The robot dropped one of the targets and was not able to pick the remaining two. The system dropped a non-target item during manipulation. The system also successfully moved 5 items between the shelf’s bins to clear occlusions for required picks. Two of those move operations allowed it two successfully pick two target products. Team Delft called the end of the run after 14 min and 45 s.

5.1 Analysis

Team Delft’s robot reliable and performing capabilities were the key to its success. Its gripper could grasp all 39 items in the competition, including the dumbbell and the pencil cup using the pinch grasp, in any orientation and bin location. However, the grasp on heavy and big items was not completely reliable. The dumbbell proved specially difficult, since the grasp approach needed extreme accuracy to succeed.

The object recognition module had an specially outstanding performance robust to varying light conditions. However, pose estimation was strongly affected by reflections, which produced scarce PointCloud data.

Most difficulties for our system were encountered when trying to find a collision-free motion plan to approach the target object. This rejected many targets that were retried later. In the next attempt, removing occluding items was done, but sometimes the cluttering of the bin caused a stall situation in which items were preventing each other from being picked.

Overall, the Picking Task proved far more difficult than the Stowing Task, with many teams scoring half the points. This is because picking from the shelf

required more manipulation, with items occluding each other. The Stowing task was basically a standard bin picking problem: all items in the tote were to be picked, and gravity helps having some easily accessible at the top. Also, the stowing in the shelf could be done with pre-computed motions to shove the target item in the bin, blindly pushing back any previous content.

5.2 Lessons Learned

Considering the results described in the previous section and the complete experience developing the robot for the Amazon Picking Challenge, we reached several conclusions about our concept design premises and how to improve it.

The most important idea is that manipulation requires contact with the environment. Team Delft's pure planning approach to grasping and manipulation treated contact as collisions to avoid, and simply by-passed this constraint for the target object. This caused a lot of rejected plans to grasp items from cluttered bins, some of them becoming actually unrealisable. Force-feedback and compliance in the gripper seem unavoidable to achieve a reliable solution. Also, creating a single gripper capable of handling such a variety of products proved difficult. None of the teams managed to pick the dumbbell, for example. Having different grippers and switching between them on the fly seems a more efficient and robust solution.

On the perception side, Deep Learning neural networks proved an excellent solution for object recognition, but they also are a really promising solution to pose estimation and even grasp planning, as the results of other teams suggest.

Notwithstanding the discussed improvements, Team Delft's concept based on speed and reliability proved successful. The ready-for-industry approach we took, with installation and setup procedures, and professional team coordination during the competition, allowed to keep robustly improving the robot's performance till reaching close to its top limit right at the competition.

6 Concluding Remarks

This paper provides a comprehensive overview of Team Delft's robot winner of the Amazon Picking Challenge 2016. The key to Delft's robot success was a concept aimed for robustness and speed, relying on an end-to-end engineering process integrating well established industry practices and cutting-edge AI and robotics technologies.

There was a new Stowing Task in the 2016 edition of the challenge, to bin-pick products from a tote and stow them in a shelf. The overall high scores by many teams, and the excellent performance of Team Delft's robot, suggest that the bin picking problem for diverse, medium-size products can be addressed by current robotic technology. Speed is still far from human performance (~ 100 items an hour, compared to 400 items an hour in the case of a human), but considering that Team Delft's robot could have been speed-up probably 50% with faster motions and faster processing, we are confident to predict that robot

technology is getting there. However, the Picking task results, proved that general manipulation, including diverse objects and cluttered spaces, still remains an open problem for robotics.

Acknowledgements. All authors gratefully acknowledge the financial support by the European Union's Seventh Framework Programme project Factory-in-a-day (FP7-609206) We would like to thank RoboValley (<http://www.robovalley.com>), the ROS-Industrial consortium, our sponsors Yaskawa, IDS, Phaer, Ikbenstil and Induvac, the people at the Delft Center for Systems and Control and TU Delft Logistics for their support, also Lacquey B.V. for helping us handle our heavy rail, and finally special thanks to Gijs vd. Hoorn for his help during the development of the robotic system.

References

1. Beeson, P., Ames, B.: TRAC-IK: an open-source library for improved solving of generic inverse kinematics. In: Proceedings of the IEEE RAS Humanoids Conference, Seoul, Korea, November 2015
2. Bohren, J.: SMACH (2016). <http://wiki.ros.org/smach>
3. Correll, N., Bekris, K.E., Berenson, D., Brock, O., Causo, A., Hauser, K., Okada, K., Rodriguez, A., Romano, J.M., Wurman, P.R.: Lessons from the amazon picking challenge. CoRR abs/1601.05484 (2016)
4. Delft Robotics, B.V. <http://www.delftrobotics.com/>
5. Factory-in-a-day. <http://www.factory-in-a-day.eu>
6. Mellado, N., Aiger, D., Mitra, N.J.: Super 4PCS fast global pointcloud registration via smart indexing. Comput. Graph. Forum **33**(5), 205–215 (2014)
7. Quigley, M., Conley, K., Gerkey, B.P., Faust, J., Foote, T., Leibs, J., Wheeler, R., Ng, A.Y.: ROS: an open-source robot operating system. In: ICRA Workshop on Open Source Software (2009)
8. Ren, S., He, K., Girshick, R., Sun, J.: Faster R-CNN: towards real-time object detection with region proposal networks. In: Advances in Neural Information Processing Systems (NIPS) (2015)
9. ROS-Industrial. <http://rosindustrial.org/>
10. Sucan, I.A., Chitta, S.: MoveIt!. <http://moveit.ros.org>
11. TUD Robotics Institute. <http://robotics.tudelft.nl>